

浙江大学

本科实验报告

课程名称：操作系统

实验名称：RV64 内核引导与时钟中断处理

姓 名：蔡雨谦

学 院：计算机学院

系：计算机系

专 业：计算机科学与技术

学 号：3210102466

指导教师：寿黎但

2023 年 9 月 28 日

浙江大学实验报告

一、实验目的

- 学习 RISC-V 汇编，编写 head.S 实现跳转到内核运行的第一个 C 函数。
- 学习 OpenSBI，理解 OpenSBI 在实验中所起到的作用，并调用 OpenSBI 提供的接口完成字符的输出。
- 学习 Makefile 相关知识，补充项目中的 Makefile 文件，来完成对整个工程的管理。
- 学习 RISC-V 的 trap 处理相关寄存器与指令，完成对 trap 处理的初始化。
- 理解 CPU 上下文切换机制，并正确实现上下文切换功能。
- 编写 trap 处理函数，完成对特定 trap 的处理。
- 调用 OpenSBI 提供的接口，完成对时钟中断事件的设置。

二、实验环境

- Ubuntu 22.04.2 LTS Windows Subsystem for Linux 2

三、操作方法与实验步骤

1. 准备工程

从课程提供的 github 仓库获取实验代码框架，准备进行后续实验。

2. RV64 内核引导

实现 RV64 内核引导需要完善的文件：

- arch/riscv/kernel/head.S
- lib/Makefile
- arch/riscv/kernel/sbi.c
- arch/riscv/include/defs.h

2.1 编写 head.S

head.S 的作用是引导内核的启动。设置一个大小为 4kb (4096) 的栈，并通过跳转指令跳转到 start_kernel 函数。

```
arch > riscv > kernel > cat head.S
1  .extern start_kernel
2
3  .section .text.entry
4  .globl _start
5  _start:
6  # -----
7  la sp, boot_stack_top
8  jal x0, start_kernel
9  # -----
10
11 .section .bss.stack
12 .globl boot_stack
13 boot_stack:
14 .space 4096 # <-- change to your stack size
15
16 .globl boot_stack_top
17 boot_stack_top:
```

2.2 完善 makefile 脚本

参照已经给出的其他 makefile 文件，将目录下的全部 .c 文件编译成 .o 文件。

```
lib > M Makefile
1  C_SRC      = $(sort $(wildcard *.c))
2  OBJ        = $(patsubst %.c,%.o,$(C_SRC))
3
4  file = printk.o
5  all: $(OBJ)
6
7  %.o: %.c
8      | ${GCC} ${CFLAG} -c $<
9  clean:
10     | $(shell rm *.o 2>/dev/null)
11
```

运行 make all 指令可以进行编译。

```
icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$ make all
make -C lib all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/lib'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fcommon -nostdinc -static -lgcc -WL,--nmagic -WL,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c printk.c
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/lib'
make -C init all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/init'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fcommon -nostdinc -static -lgcc -WL,--nmagic -WL,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c main.c
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fcommon -nostdinc -static -lgcc -WL,--nmagic -WL,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c test.c
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/init'
make -C arch/riscv all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/arch/riscv'
make -C kernel all
make[2]: Entering directory '/home/icey/os23fall-stu/src/lab1/arch/riscv/kernel'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fcommon -nostdinc -static -lgcc -WL,--nmagic -WL,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c head.S
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fcommon -nostdinc -static -lgcc -WL,--nmagic -WL,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c sbi.c
make[2]: Leaving directory '/home/icey/os23fall-stu/src/lab1/arch/riscv/kernel'
riscv64-linux-gnu-ld -T kernel/vmlinux.lds kernel/*.o ../init/*.o ../lib/*.o -o ../vmlinux
riscv64-linux-gnu-objcopy -O binary ../vmlinux ./boot/Image
rm ../vmlinux > ../System.map
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/arch/riscv'

Build Finished OK
icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$
```

编译后会生成 vmlinux 内核文件，使用 file vmlinux 指令查看：

```
icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$ file vmlinux
vmlinux: ELF 64-bit LSB executable, UCB RISC-V, soft-float ABI, version 1 (SYSV), statically linked, with debug_info, not stripped
icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$
```

2.3 补充 sbi.c

参照示例一，使用 asm volatile 向 C 代码中插入内联汇编代码。将 ext 放入寄存器 a7 中，fid 放入寄存器 a6 中，将 arg0 - arg5 放入寄存器 a0 - a5 中。使用 ecall 指令。ecall 之后系统会进入 M 模式。返回结果存放在寄存器 a0、a1 中，其中 a0 为 error，a1 为返回值。将 a0、a1 的值放进一个 sbiret 变量中作为 sbi_ecall 的返回值。

```

3
4 struct sbiret sbi_ecall(int ext, int fid, uint64 arg0,
5                       uint64 arg1, uint64 arg2,
6                       uint64 arg3, uint64 arg4,
7                       uint64 arg5)
8 {
9     // unimplemented
10    struct sbiret ret;
11    __asm__ volatile(
12        "mv a7, %[ext]\n"
13        "mv a6, %[fid]\n"
14        "mv a0, %[arg0]\n"
15        "mv a1, %[arg1]\n"
16        "mv a2, %[arg2]\n"
17        "mv a3, %[arg3]\n"
18        "mv a4, %[arg4]\n"
19        "mv a5, %[arg5]\n"
20        "ecall\n"
21        "mv %[error], a0\n"
22        "mv %[value], a1\n"
23        : [error] "=r"(ret.error), [value] "=r"(ret.value)
24        : [ext] "r"(ext), [fid] "r"(fid), [arg0] "r"(arg0), [arg1] "r"(arg1),
25        [arg2] "r"(arg2), [arg3] "r"(arg3), [arg4] "r"(arg4), [arg5] "r"(arg5)
26        : "a0", "a1", "a2", "a3", "a4", "a5", "a6", "a7");
27    return ret;
28 }
29

```

2.4 修改 defs.h

参照示例二，使用 `asm volatile` 插入内联汇编代码，完善 `csr_read` 宏，将 `csr` 寄存器的值读到变量 `_v` 中。

```

6  #define csr_read(csr) \
7  ({ \
8      register uint64 __v; \
9      /* unimplemented */ \
10     asm volatile ( \
11         "csrr %[v], " #csr \
12         : [v] "=r" (__v) \
13         : \
14         : "memory" \
15     ); \
16     __v; \
17 })
18

```

代码框架中给出的 `csr_write` 宏则是将 `val` 的值写进 `csr` 寄存器中。

```

19 #define csr_write(csr, val) \
20 ({ \
21     uint64 __v = (uint64)(val); \
22     asm volatile ("csrw " #csr ", %0" \
23         : : "r" (__v) \
24         : "memory"); \
25 })
26
27 #endif
28

```

2.5 运行 RV64 内核。

使用 `make all` 指令编译项目文件夹得到的内核可以通过 `make`

run 指令运行。

[illegible]

3. RV64 时钟中断处理

实现 RV64 时钟中断处理需要完善或添加的文件:

- arch/riscv/kernel/head.S
- arch/riscv/kernel/entry.S
- arch/riscv/kernel/trap.c
- arch/riscv/kernel/clock.c

3.1 准备工作

vmlinux.lds 与 head.S 文件中有需要修改的部分如下(左为修改前, 右为修改后):

<pre> 20 .text : ALIGN(0x1000){ 21 _stext = .; 22 23 *(&.text.entry) 24 *(&.text .text.*) 25 26 _etext = .; 27 } 28 </pre>	<pre> 20 .text : ALIGN(0x1000){ 21 _stext = .; 22 23 *(&.text.init) 24 *(&.text.entry) 25 *(&.text .text.*) 26 27 _etext = .; 28 } 29 </pre>
--	---

<pre> 3 .section .text.entry 4 .globl _start 5 _start: 6 # ----- </pre>	<pre> 3 .section .text.init 4 .globl _start 5 _start: 6 # ----- </pre>
---	--

3.2 开启 trap 处理

在运行 start_kernel 前需要在 head.S 中对一些寄存器进行初始化:

设置 stvec, 将 _traps 所表示的地址写入 stvec

```

# set stvec = _traps
la a0, _traps
csrw stvec, a0

```

开启时钟中断, 将 sie[STIE]置 1, STIE 是第五位:

```

# set sie[STIE] = 1
li a0, 1<<5
csrw sie, a0

```

设置第一次时钟中断

```

# set first time interrupt
jal ra, clock_set_next_event

```

开启 S 态下的中断响应, 将 sstatus[SIE]置 1, SIE 是第一位:

```

# set sstatus[SIE] = 1
li a0, 1<<1
csrw sstatus, a0

```

3.3 实现上下文切换

在 arch/riscv/kernel/ 目录下添加 entry.S 文件, 文件包含以下内容:

保存 CPU 的寄存器到内存中

```

# 1. save 32 registers and sepc to stack
addi sp, sp, -256
sd x1, 0(sp)
sd x2, 8(sp)
sd x3, 16(sp)
sd x4, 24(sp)
sd x5, 32(sp)
sd x6, 40(sp)
sd x7, 48(sp)
sd x8, 56(sp)
sd x9, 64(sp)
sd x10, 72(sp)
sd x11, 80(sp)
sd x12, 88(sp)
sd x13, 96(sp)
sd x14, 104(sp)
sd x15, 112(sp)
sd x16, 120(sp)
sd x17, 128(sp)
sd x18, 136(sp)
sd x19, 144(sp)
sd x20, 152(sp)
sd x21, 160(sp)
sd x22, 168(sp)
sd x23, 176(sp)
sd x24, 184(sp)
sd x25, 192(sp)
sd x26, 200(sp)
sd x27, 208(sp)
sd x28, 216(sp)
sd x29, 224(sp)
sd x30, 232(sp)
sd x31, 240(sp)
csrr t0, sepc
sd t0, 248(sp)
-----

```

将 scause 和 sepc 中的值传入 trap 处理函数 trap_handler

```

# 2. call trap_handler
csrr a0, scause
csrr a1, sepc
jal x1, trap_handler

```

从内存中恢复 CPU 的寄存器

```

# 3. restore sepc and 32 registers (x2(sp) should be restored)
ld t0, 248(sp)
csrw sepc, t0
ld x1, 0(sp)
ld x3, 16(sp)
ld x4, 24(sp)
ld x5, 32(sp)
ld x6, 40(sp)
ld x7, 48(sp)
ld x8, 56(sp)
ld x9, 64(sp)
ld x10, 72(sp)
ld x11, 80(sp)
ld x12, 88(sp)
ld x13, 96(sp)
ld x14, 104(sp)
ld x15, 112(sp)
ld x16, 120(sp)
ld x17, 128(sp)
ld x18, 136(sp)
ld x19, 144(sp)
ld x20, 152(sp)
ld x21, 160(sp)
ld x22, 168(sp)
ld x23, 176(sp)
ld x24, 184(sp)
ld x25, 192(sp)
ld x26, 200(sp)
ld x27, 208(sp)
ld x28, 216(sp)
ld x29, 224(sp)
ld x30, 232(sp)
ld x31, 240(sp)
ld x2, 8(sp)
addi sp, sp, 256

```

从 trap 中返回

```

# 4. return from trap
sret

```

3.4 实现 trap 处理函数

在 arch/riscv/kernel/ 目录下添加 trap.c 文件，文件包含 trap 处理函数 trap_handler()，scause 和 sepc 寄存器的值作为参数。scause 寄存器最高位为 1 时为 interrupt，剩余位表示 5 时为 timer interrupt。在 arch/riscv/include/ 目录下添加 对应 trap.h 文件


```

5 void trap_handler(unsigned long scause, unsigned long sepc) {
6     // 通过 `scause` 判断trap类型
7     // 如果是interrupt 判断是否是timer interrupt
8     // 如果是timer interrupt 则打印输出相关信息, 并通过 `clock_set_next_event()` 设置下一次时钟中断
9     // `clock_set_next_event()` 见 4.3.4 节
10    // 其他interrupt / exception 可以直接忽略
11
12    // YOUR CODE HERE
13    if (scause >> 63){ // 最高位为1就是interrupt
14        unsigned long rest = scause - 0x8000000000000000;
15        if (rest == 5){ // 剩余位表示5时是timer interrupt
16            printk("[S] Supervisor Mode Timer Interrupt\n");
17            clock_set_next_event();
18        }
19    }
20 }

```

```

arch > riscv > include > C trap.h > ...
1  #ifndef _TRAP_H
2  #define _TRAP_H
3
4  void trap_handler(unsigned long scause, unsigned long sepc);
5  #endif

```

3.5 实现时钟中断相关函数

在 arch/riscv/kernel/ 目录下添加 clock.c 文件。文件包含两个函数：get_cycles()：使用 rdttime 汇编指令获得当前 time 寄存器中的值；clock_set_next_event()：调用 sbi_ecall，设置下一个时钟中断事件。

```

unsigned long get_cycles() {
    // 编写内联汇编, 使用 rdttime 获取 time 寄存器中 (也就是mtime 寄存器) 的值并返回
    // YOUR CODE HERE
    unsigned long ret;
    __asm__ volatile (
        "rdtime %[ret]\n"
        : [ret] "=r" (ret)
        :
        : "memory"
    );
    return ret;
}

```

```

void clock_set_next_event() {
    // 下一次 时钟中断 的时间点
    unsigned long next = get_cycles() + TIMECLOCK;

    // 使用 sbi_ecall 来完成对下一次时钟中断的设置
    // YOUR CODE HERE
    sbi_ecall(0,0,next,0,0,0,0,0);
}

```

3.6 编译测试

虽然新添了.c 文件, 但是在 RV64 内核引导部分写的 makefile

文件将全部.c文件一起处理了，所以不用修改。

```
icey@LAPTOP-RHUCAGPI:~/os23fall-stu/src/lab1$ make all
make -C lib all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/lib'
make[1]: Nothing to be done for 'all'.
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/lib'
make -C init all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/init'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c main.c
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c test.c
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/init'
make -C arch/riscv all
make[1]: Entering directory '/home/icey/os23fall-stu/src/lab1/arch/riscv'
make -C kernel all
make[2]: Entering directory '/home/icey/os23fall-stu/src/lab1/arch/riscv/kernel'
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c entry.S
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c head.S
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c clock.c
riscv64-linux-gnu-gcc -march=rv64imafd -mabi=lp64 -mcmodel=medany -fno-builtin -ffunction-sections -fdata-sections -nostartfiles -nostdlib -nostdinc -static -lgcc -Wl,--nmagic -Wl,--gc-sections -g -I /home/icey/os23fall-stu/src/lab1/include -I /home/icey/os23fall-stu/src/lab1/arch/riscv/include -c trap.c
make[2]: Leaving directory '/home/icey/os23fall-stu/src/lab1/arch/riscv/kernel'
riscv64-linux-gnu-ld -T kernel/vmlinux.lds kernel/*.o ../init/*.o ../lib/*.o -o ../vmlinux
riscv64-linux-gnu-objcopy -O binary ../vmlinux ./boot/image
nm ../vmlinux > ../System.map
make[1]: Leaving directory '/home/icey/os23fall-stu/src/lab1/arch/riscv'
Build Finished OK
```

运行 make run 指令，有如下输出：

```
Boot HART PMP Granularity : 4
Boot HART PMP Address Bits: 54
Boot HART MHPM Count : 0
Boot HART MHPM Count : 0
Boot HART MIDELEG : 0x00000000000000222
Boot HART MEDELEG : 0x0000000000000b109
2022 Hello RISC-V
ssstatus: 2
sscratch: 0
sscratch after write: 9
kernel is running!
[S] Supervisor Mode Timer Interrupt
kernel is running!
[S] Supervisor Mode Timer Interrupt
kernel is running!
[S] Supervisor Mode Timer Interrupt
kernel is running!
[S] Supervisor Mode Timer Interrupt
kernel is running!
[S] Supervisor Mode Timer Interrupt
kernel is running!
[S] Supervisor Mode Timer Interrupt
```

四、思考题

1. 请总结一下 RISC-V 的 calling convention，并解释 Caller / Callee Saved Register 有什么区别？

calling convention(函数调用规范):

- 1) 将参数存储到函数能够访问到的位置;
- 2) 跳转到函数开始位置 (使用 jal 指令)
- 3) 获取函数需要的局部存储资源，按需保留寄存器
- 4) 执行函数中的指令
- 5) 将返回值存储到调用者能够访问到的位置，回复寄存器，释放局部存储资源
- 6) 返回调用函数的位置 (使用 ret 指令)

Caller / Callee Saved Register 的区别:

caller saved registers (调用者保存寄存器) 也叫易失性寄存器，在程序调用的过程中，这些寄存器中的值不需要被保存，如果某一个程序需要保存这个寄存器的值，需要调用者自己压入栈;

callee saved registers (被调用者保存寄存器) 也叫非易失性寄存器，在程序调用过程中，

这些寄存器中的值需要被保存，不能被覆盖；当某个程序调用这些寄存器，被调用寄存器会先保存这些值然后再进行调用，且在调用结束后恢复被调用之前的值；

2. 编译之后，通过 System.map 查看 vmlinux.lds 中自定义符号的值（截图）。

```
icey@LAPTOP-RKUCA6PT:~/os23fall-stu/src/lab1$ nm vmlinux
0000000080200000 t $x
000000008020000c t $x
00000000802000f0 t $x
0000000080200134 t $x
0000000080200144 t $x
0000000080200194 t $x
0000000080200670 t $x
0000000080200000 A _BASE_ADDR
0000000080202000 d _GLOBAL_OFFSET_TABLE_
0000000080204000 B _ebss
0000000080202000 D _edata
0000000080204000 B _ekernel
0000000080201070 R _erodata
00000000802006f0 T _etext
0000000080203000 B _sbss
0000000080202000 D _sdata
0000000080200000 T _skernel
0000000080201000 R _srodata
0000000080200000 T _start
0000000080200000 T _stext
0000000080203000 B boot_stack
0000000080204000 B boot_stack_top
0000000080200670 T printk
0000000080200144 T putc
000000008020000c T sbi_ecall
00000000802000f0 T start_kernel
0000000080200134 T test
0000000080200194 t vprintfmt
icey@LAPTOP-RKUCA6PT:~/os23fall-stu/src/lab1$
```

3. 用 csr_read 宏读取 sstatus 寄存器的值，对照 RISC-V 手册解释其含义（截图）。

4. 用 csr_write 宏向 sscratch 寄存器写入数据，并验证是否写入成功（截图）。

3 and 4:

要输出寄存器的值，需要一个输出函数，负责输出 int 类型的值：

```
115 void printi(int x)
116 {
117     char X[16];
118     int i = 0;
119     if (x == 0) {
120         putc('0');
121         return;
122     }
123
124     while (x) {
125         X[i++] = '0' + x % 10;
126         x /= 10;
127     }
128
129     while (i) {
130         putc(X[--i]);
131     }
132 }
```

在 main.c 中添加代码用来读写寄存器：

```
10 register int val;
11 val = csr_read(sstatus);
12 printk("sstatus: ");
13 printi(val);
14 printk("\n");
15
16 val = csr_read(sscratch);
17 printk("sscratch: ");
18 printi(val);
19 printk("\n");
20 csr_write(sscratch,9);
21 val = csr_read(sscratch);
22 printk("sscratch after write: ");
23 printi(val);
24 printk("\n");
25
```

输出如下：

```

Boot HART MEDELEG      : 0x000000000000000b
2022 Hello RISC-V
sstatus: 2
sscratch: 0
sscratch after write: 9
[6] Supervisor Mode Timer Test complete

```

sstatus 寄存器的第 1 位(SIE)为 1 说明允许 S 模式的中断。
sscratch 寄存器前后值的变化说明写入成功。

5. Detail your steps about how to get arch/arm64/kernel/sys.i

进入 linux 文件夹：

```

icey@LAPTOP-RKUCA6PI:/mnt/d/文档/course/OS/lab/lab0$ cd linux-6.6-rc2/

```

配置 defconfig:

```

icey@LAPTOP-RKUCA6PI:/mnt/d/文档/course/OS/lab/lab0/linux-6.6-rc2$ make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- defconfig
*** Default configuration is based on 'defconfig'
#
# configuration written to .config
#

```

生成 sys.i 文件：

```

#
icey@LAPTOP-RKUCA6PI:/mnt/d/文档/course/OS/lab/lab0/linux-6.6-rc2$ make ARCH=arm64 CROSS_COMPILE=aarch64-linux-gnu- arch/arm64/kernel/sys.i
SYNC    include/config/auto.conf.cmd
WRAP    arch/arm64/include/generated/uapi/asm/kvm_para.h
WRAP    arch/arm64/include/generated/uapi/asm/errno.h
WRAP    arch/arm64/include/generated/uapi/asm/ioctl.h
WRAP    arch/arm64/include/generated/uapi/asm/ioctls.h
WRAP    arch/arm64/include/generated/uapi/asm/inchuf.h

```

```

LD      arch/arm64/kernel/vdso/vdso.so.dbg
VDSOSYM include/generated/vdso-offsets.h
OBJCOPY arch/arm64/kernel/vdso/vdso.so
CPP     arch/arm64/kernel/sys.i

```

6. Find system call table of Linux v6.0 for ARM32, RISC-V(32 bit), RISC-V(64 bit), x86(32 bit), x86_64 List source code file, the whole system call table with macro expanded, screenshot every step.

System call table for x86:

document > course > OS > lab > lab0 > linux-6.6-rc2 > arch > x86 > entry > syscalls				
<div> <div> <div></div> <div></div> <div></div> <div></div> </div> <div>Sort</div> <div>View</div> <div></div> </div>				
<input type="checkbox"/> 名称	修改日期	类型	大小	
Makefile	2023/9/18 5:40	文件	3 KB	
syscall_32.tbl	2023/9/18 5:40	TBL 文件	18 KB	
syscall_64.tbl	2023/9/18 5:40	TBL 文件	15 KB	

如图所示：

```

#
# 32-bit system call numbers and entry vectors
#
# The format is:
# <number> <abi> <name> <entry point> <compat entry point>
#
# The __ia32_sys and __ia32_compat_sys stubs are created on-the-fly for
# sys.*() system calls and compat_sys.*() compat system calls if
# IA32_EMULATION is defined, and expect struct pt_regs *regs as their only
# parameter.
#
# The abi is always "i386" for this file.
#
0   i386  restart_syscall      sys_restart_syscall
1   i386  exit                 sys_exit
2   i386  fork                 sys_fork
3   i386  read                 sys_read
4   i386  write                sys_write
5   i386  open                 sys_open          compat_sys_open
6   i386  close                sys_close
7   i386  waitpid              sys_waitpid
8   i386  creat                sys_creat
9   i386  link                 sys_link
10  i386  unlink               sys_unlink
11  i386  execve               sys_execve
    compat_sys_execve
12  i386  chdir                sys_chdir
13  i386  time                 sys_time32
14  i386  mknod                sys_mknod
15  i386  chmod                sys_chmod
16  i386  lchown               sys_lchown16
17  i386  break
18  i386  oldstat              sys_stat
19  i386  lseek                sys_lseek          compat_sys_lseek
20  i386  getpid               sys_getpid
21  i386  mount                sys_mount
22  i386  umount               sys_oldumount
23  i386  setuid               sys_setuid16
    ..
    ..
    ..

```

7. Explain what is ELF file? Try readelf and objdump command on an ELF file, give screenshot of the output. Run an ELF file and cat /proc/PID/maps to give its memory layout.

ELF(executable and linkable format)是一种可执行文件格式，主要用于 linux 平台，windows 平台下相对应的文件为 coff 格式。

```

icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$ readelf -a vmlinux
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                           ELF64
  Data:                               2's complement, little endian
  Version:                           1 (current)
  OS/ABI:                            UNIX - System V
  ABI Version:                        0
  Type:                              EXEC (Executable file)
  Machine:                           RISC-V
  Version:                           0x1
  Entry point address:                0x80200000
  Start of program headers:          64 (bytes into file)
  Start of section headers:         22280 (bytes into file)
  Flags:                              0x0
  Size of this header:                64 (bytes)
  Size of program headers:           56 (bytes)
  Number of program headers:          4
  Size of section headers:           64 (bytes)
  Number of section headers:         20
  Section header string table index: 19

```

```

Section Headers:
 [Nr] Name              Type              Address             Offset
     Size              EntSize          Flags Link    Info  Align
 [ 0] 0000000000000000 NULL              0000000000000000  0 0 0
 [ 1] .text               PROGBITS          0000000000000000  0 0 4096
     0000000000000934 0000000000000000 AX      0 0
 [ 2] .rodata             PROGBITS          0000000000000000  0 0 4096
     0000000000000098 0000000000000000 A       0 0
 [ 3] .data               PROGBITS          0000000000000000  0 0 4096
     0000000000000008 0000000000000000 WA      0 0
 [ 4] .got                PROGBITS          0000000000000008  0 0 8
     0000000000000018 0000000000000008 WA      0 0
 [ 5] .got.plt            PROGBITS          0000000000000020  0 0 8
     0000000000000010 0000000000000008 WA      0 0
 [ 6] .bss                NOBITS            0000000000000000  0 0 4096
     0000000000000100 0000000000000000 WA      0 0
 [ 7] .debug_info         PROGBITS          0000000000000000  0 0 3072

```



```

Key to Flags:
W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
L (link order), O (extra OS processing required), G (group), T (TLS),
C (compressed), x (unknown), o (OS specific), E (exclude),
D (mbind), p (processor specific)

There are no section groups in this file.

Program Headers:
Type           Offset             VirtAddr           PhysAddr
               FileSiz            MemSiz             Flags             Align
RISCV_ATTRIBUT 0x0000000000004cc5 0x0000000000000000 0x0000000000000000
               0x000000000000032 0x0000000000000000 R                 0x1
LOAD           0x0000000000001000 0x0000000000020000 0x0000000000020000
               0x0000000000001098 0x0000000000001098 R E              0x1000
LOAD           0x0000000000003000 0x0000000000020200 0x0000000000020200
               0x000000000000030 0x0000000000002000 RW              0x1000
GNU_STACK      0x0000000000000000 0x0000000000000000 0x0000000000000000
               0x0000000000000000 0x0000000000000000 RW              0x10

Section to Segment mapping:
Segment Sections...
00      .riscv.attributes
01      .text .rodata
02      .data .got .got.plt .bss
03

```

```

Symbol table '.symtab' contains 63 entries:
Num:   Value              Size Type      Bind   Vis      Ndx Name
0: 0000000000000000      0 NOTYPE    LOCAL DEFAULT UND
1: 0000000000020000      0 SECTION   LOCAL DEFAULT 1 .text
2: 0000000000020100      0 SECTION   LOCAL DEFAULT 2 .rodata
3: 0000000000020200      0 SECTION   LOCAL DEFAULT 3 .data
4: 0000000000020208      0 SECTION   LOCAL DEFAULT 4 .got
5: 0000000000020220      0 SECTION   LOCAL DEFAULT 5 .got.plt
6: 0000000000020300      0 SECTION   LOCAL DEFAULT 6 .bss
7: 0000000000000000      0 SECTION   LOCAL DEFAULT 7 .debug_info
8: 0000000000000000      0 SECTION   LOCAL DEFAULT 8 .debug_abbrev
9: 0000000000000000      0 SECTION   LOCAL DEFAULT 9 .debug_aranges
10: 0000000000000000      0 SECTION   LOCAL DEFAULT 10 .debug_rnglists
11: 0000000000000000      0 SECTION   LOCAL DEFAULT 11 .debug_line
12: 0000000000000000      0 SECTION   LOCAL DEFAULT 12 .debug_str
13: 0000000000000000      0 SECTION   LOCAL DEFAULT 13 .debug_line_str
14: 0000000000000000      0 SECTION   LOCAL DEFAULT 14 .comment
15: 0000000000000000      0 SECTION   LOCAL DEFAULT 15 .riscv.attributes
16: 0000000000000000      0 SECTION   LOCAL DEFAULT 16 .debug_frame
17: 0000000000000000      0 FILE      LOCAL DEFAULT ABS head.o
18: 0000000000020000      0 NOTYPE    LOCAL DEFAULT 1 $x
19: 0000000000000000      0 FILE      LOCAL DEFAULT ABS entry.o
20: 0000000000020028      0 NOTYPE    LOCAL DEFAULT 1 $x
21: 0000000000000000      0 FILE      LOCAL DEFAULT ABS clock.c
22: 0000000000020148      0 NOTYPE    LOCAL DEFAULT 1 $x

```

```

icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$ objdump -v vmlinux
GNU objdump (GNU Binutils for Ubuntu) 2.38
Copyright (C) 2022 Free Software Foundation, Inc.
This program is free software; you may redistribute it under the terms of
the GNU General Public License version 3 or (at your option) any later version.
This program has absolutely no warranty.

```

```

-bash: ./System.map: Permission denied
icey@LAPTOP-RKUCA6PI:~/os23fall-stu/src/lab1$ cat System.map
0000000000000000 t $x
000000000000002c t $x
0000000000000014c t $x
00000000000000174 t $x
000000000000001d8 t $x
000000000000002a4 t $x
00000000000000308 t $x
000000000000003f4 t $x
00000000000000404 t $x
00000000000000454 t $x
00000000000000930 t $x
000000000000009b0 t $x
0000000000020000 A _BASE_ADDR
0000000000020200 D _TIMECLOCK
0000000000020208 d _GLOBAL_OFFSET_TABLE_
0000000000020400 B _ebss
0000000000020908 D _edata
0000000000020400 B _kernel
000000000002010e8 R _erodata
00000000000200a80 T _etext
00000000000203000 B _sbss
0000000000020200 D _sdata
0000000000020000 T _skernel
00000000000201000 R _srodata
0000000000020000 T _start
0000000000020000 T _stext
000000000002002c T _traps
00000000000203000 B boot_stack
00000000000204000 B boot_stack_top
00000000000200174 T clock_set_next_event
0000000000020014c T get_cycles

```

8. 在我们使用 `make run` 时，OpenSBI 会产生如下输出：

4、实现时钟中断的实验中，在往 head.S 文件增添代码时，要添在加载栈地址的代码之后。

```

5  _start:
6
7      # 内核引导部分的代码
8
9      la sp, boot_stack_top
10
11     # 时钟中断部分的代码
12     # -----
13
14     # set stvec = _traps
15     la a0, _traps
16     csrw stvec, a0
17
18     # -----
19
20     # set sie[STIE] = 1
21     li a0, 1<<5
22     csrw sie, a0

```